
Supplementary Material: Pedagogical State Tracking: Extending Knowledge Tracing to Whole-Classroom Discourse

This Supplementary Material covers every model from first principles:
motivation, mathematical formulation, prior choices,
forward algorithm derivation, worked numerical examples,
and code. Updated version of this Supplementary Material can be found on
https://github.com/drake1729/AIML_System

Contents

1	What Are We Modelling?	3
1.1	The Problem	3
1.2	Data at a Glance	3
1.3	The Covariate Vector	3
1.4	The Three Models and Their Trade-Offs	4
2	Model 1: BLFM — Bayesian Latent Factor Model	5
2.1	Core Question	5
2.2	The Exchangeability Assumption	5
2.3	The Sufficient Statistic: Episode-Level Aggregation	5
2.4	Model Specification	6
2.4.1	Step 1 — Automatic Relevance Determination (ARD) Prior	6
2.4.2	Step 2 — Intercept and Random Effects	6
2.4.3	Step 3 — Predicted Episode Rate	6
2.4.4	Step 4 — Beta Regression Likelihood	7
2.5	Full Generative Story	7
2.6	Worked Numerical Example	7
2.7	Code	8
2.7.1	Key Implementation Details Explained	9
2.8	Running MCMC	9
2.9	What BLFM Cannot Answer	10
3	Model 2: PST — Probabilistic Skill Tracing (Sequential BKT)	10
3.1	Core Question	10
3.2	Hidden Markov Model (HMM) Primer	10
3.3	PST Parameters	10
3.3.1	The Transition Matrix	11
3.3.2	Emission Probabilities	11
3.3.3	Why Are p_S and p_G Fixed?	12
3.4	The Forward Algorithm: Step-by-Step Derivation	12
3.4.1	Why Not Enumerate All State Sequences?	12
3.4.2	Notation	12
3.4.3	Initialisation ($t = 1$)	12
3.4.4	Recursion ($t \geq 2$): Predict–Update	13
3.4.5	Log-Likelihood	13
3.5	Worked Numerical Forward Pass	13
3.6	Prior Choices Explained	14
3.7	Key Result: Regime Persistence	14
3.8	Code	14
3.8.1	Two Technical Details	16
3.9	Why PST’s AUC is Misleading	16
4	Model 3: IO-BKT — Input-Output Bayesian Knowledge Tracing	17
4.1	Core Question	17
4.2	The Key Innovation: Content-Driven Transitions	17

4.3	Full Model Specification	18
4.3.1	Priors	18
4.3.2	Emission Model	18
4.3.3	Forward Algorithm with Time-Varying Transitions	19
4.4	Worked Numerical Example: Transition Probabilities	19
4.5	Interpreting Table II: The Posterior Coefficients	19
4.6	Code	20
4.6.1	The <code>einsum</code> Operation Decoded	22
4.6.2	MCMC Configuration	22
4.7	Label Switching: Detection and Correction	23
4.8	The Full-Length Normalisation (Tmax Extension)	24
5	Experimental Setup	25
6	Results and Discussions	25
6.1	Baselines	25
6.2	What Each AUC Gap Means	25
6.3	PST: Regime Persistence Finding	26
6.4	The Three-Stage AI Coaching Pipeline	27
6.5	Summary: The Three Models Side by Side	28
6.6	Conclusion	28
6.7	Knowledge Change Curves	29
6.7.1	Motivation	29
6.7.2	Notation	29
6.7.3	Episode Opening Baseline	29
6.7.4	Feature Occurrence Indexing	30
6.7.5	Per-Occurrence Change from Opening	30
6.7.6	The Knowledge Change Curve	30
6.7.7	Axis Definitions	31
6.7.8	Why $P(M_t=1)$ and Not O_t	31
6.7.9	Comparison with Scarlatos et al.	31
6.7.10	Interpretation of Fig. 3	31
6.8	Teacher $\rightarrow$ Student Lag Analysis	32
6.8.1	Why “Lag” Analysis	32
6.8.2	Constructing Matched Pairs	32
6.8.3	EP-Conditional Student 21CS Rate	33
6.8.4	EP-Conditional Student Cognitive Process Level	33
6.8.5	Lesson-Phase Analysis via Global Encoding	33
6.8.6	Results	34

1 What Are We Modelling?

1.1 The Problem

Teaching is sequential: a teacher speaks, students respond, the cognitive climate of the room shifts. The paper asks: *can we track whether a classroom is in a “productive discourse regime”, one that activates 21st-Century Skills (21CS), utterance by utterance?*

This is discourse Pedagogical State Tracking. It adapts Bayesian Knowledge Tracing (BKT), originally designed for tracking a single student’s mastery of a single topic, to the whole-classroom level.

1.2 Data at a Glance

Property	Value	Note
Episodes E	49 (48 usable)	31 Science, 18 Math; Grades 6–12
Total lines	19,181	Median episode: 384 lines
Mean duration	28.7 min	SD 4.8; range 12–39 min
Source	Indian classroom (48/49)	1 episode from vignette named, ‘Volcanoes? Here?’
Feature	Dimensions	Description
EP (Educational Practice)	3	Behaviourism, Cognitivism, Constructivism
KD (Knowledge Dimension)	4	Factual, Conceptual, Procedural, Metacognitive
CP (Cognitive Process)	6	Remembering, Understanding, Applying, Analysing, Evaluating, Creating
Outcome		
$O_t \in \{0, 1\}$	1	1 if any 21CS active at line t (prevalence 0.482)
21CS	4	Problem Solving, Creativity, Collaborative Learning, Critical Thinking

1.3 The Covariate Vector

At every line t , we have a 14-dimensional binary vector (13 content features + 1 intercept):

$$\mathbf{z}_t = \underbrace{[\text{Beh, Cog, Con}]}_{3 \text{ EP}} \parallel \underbrace{[\text{Fac, Con, Pro, Met}]}_{4 \text{ KD}} \parallel \underbrace{[\text{Rem, Und, App, Ana, Eva, Cre}]}_{6 \text{ CP}} \parallel \underbrace{[1]}_1 \in \{0, 1\}^{13} \times \{1\} \quad (1)$$

1.4 The Three Models and Their Trade-Offs

Model	Core Assumption	Temporal?	Content-Driven?	AUC
BLFM	Episode = bag of lines	No	Yes (mean)	0.772
PST	HMM, fixed transitions	Yes	No	0.928*
IO-BKT	HMM, content-driven trans.	Yes	Yes	0.873

* Inflated by autocorrelation, explained in §3.9.

The models are not competitors: they *answer different questions* and form a deliberate complexity ladder. We start with the simplest.

2 Model 1: BLFM — Bayesian Latent Factor Model

2.1 Core Question

BLFM Answers

Does the average mix of pedagogical moves in an episode predict how often 21CS will be active, regardless of the order in which those moves appear?

2.2 The Exchangeability Assumption

BLFM makes one fundamental assumption: lines within an episode are **exchangeable** — like balls drawn from a bag. Shuffling the lines in any order produces the same prediction.

This is, at first, startling. Why would we ignore temporal order?

Why Exchangeability Is a Reasonable Starting Point

PST (Model 2) estimates $\hat{p}_L = 0.073$: the classroom changes regime roughly *once every 14 lines*. An episode with 384 lines sees only ~ 27 transitions. The classroom spends most of its time locked in one regime. Therefore, the *fraction* of lines spent in each regime (captured by the mean feature vector) is a reasonable proxy for the dominant regime quality. Order only matters at the few transition moments.

2.3 The Sufficient Statistic: Episode-Level Aggregation

For episode e with T_e lines, BLFM collapses the full sequence into two summaries:

$$\bar{\mathbf{x}}_e = \frac{1}{T_e} \sum_{t=1}^{T_e} \mathbf{z}_t \in [0, 1]^{13} \quad (\text{mean feature vector}) \quad (2)$$

$$\bar{y}_e = \frac{1}{T_e} \sum_{t=1}^{T_e} O_t \in (0, 1) \quad (\text{observed 21CS rate}) \quad (3)$$

The 19,181 individual lines collapse to **48 episode-level observations**. BLFM regresses $\bar{y}_e$ on $\bar{\mathbf{x}}_e$.

Computing the Sufficient Statistic

Suppose Episode 7 has $T_e = 100$ lines:

- 70 lines: Behaviourism=1 $\Rightarrow \bar{x}_{\text{Beh}} = 0.70$
- 20 lines: Cognitivism=1 $\Rightarrow \bar{x}_{\text{Cog}} = 0.20$
- 10 lines: Constructivism=1 $\Rightarrow \bar{x}_{\text{Con}} = 0.10$
- 40 lines: $O_t = 1 \Rightarrow \bar{y}_7 = 0.40$

BLFM sees $(\bar{\mathbf{x}}_7, 0.40)$ as a **single data point**. It does not know which specific lines had 21CS active, nor whether the constructivist moves came at the beginning or end.

2.4 Model Specification

2.4.1 Step 1 — Automatic Relevance Determination (ARD) Prior

Instead of assigning all features the same prior variance, BLFM gives each feature f its own *precision* α_f (inverse variance):

$$\alpha_f \sim \text{Gamma}(1, 1) \quad \forall f = 1, \dots, P \quad (4)$$

$$w_f \sim \mathcal{N}\left(0, \frac{1}{\alpha_f}\right) \quad \forall f \quad (5)$$

ARD Intuition — Automatic Feature Selection

Think of α_f as a “tightness dial” for feature f .

- **Irrelevant feature:** The data provides no signal about w_f . α_f stays large (the $\text{Gamma}(1, 1)$ prior has mode 0, meaning it prefers small but the posterior is driven by likelihood). Large $\alpha_f \Rightarrow$ small variance $1/\alpha_f \Rightarrow w_f \approx 0$. The feature is effectively *switched off*.
- **Relevant feature:** The data pulls α_f toward small values, widening the prior and allowing w_f to be non-zero.

Result: you never need to manually select which of the 13 EP/KD/CP features to include. The model selects them automatically.

Implementation note. The standard deviation of w_f is $1/\sqrt{\alpha_f}$, so in NumPyro: `dist.Normal(0, 1.0/jnp.sqrt(alpha))`. The precision is α_f , not the standard deviation.

2.4.2 Step 2 — Intercept and Random Effects

$$b \sim \mathcal{N}(0, 1) \quad (6)$$

$$\sigma_{\text{cell}} \sim \text{Half-Normal}(1) \quad (7)$$

$$u_e \mid \sigma_{\text{cell}} \sim \mathcal{N}(0, \sigma_{\text{cell}}^2) \quad (8)$$

u_e is a **per-episode random effect** capturing everything the content features cannot explain: the specific teacher’s charisma, the class’s prior engagement, time of day, and so on. Episodes sharing the same Subject×Grade cell share a common σ_{cell} , creating **partial pooling**: a Grade-9 Science episode borrows strength from other Grade-9 Science episodes when estimating σ_{cell} .

2.4.3 Step 3 — Predicted Episode Rate

$$\boxed{\mu_e = \sigma(\mathbf{w}^\top \bar{\mathbf{x}}_e + u_e + b)} \quad \mu_e \in (0, 1) \quad (9)$$

where $\sigma(x) = 1/(1 + e^{-x})$ is the sigmoid/logistic function that maps any real number to $(0, 1)$.

2.4.4 Step 4 — Beta Regression Likelihood

Since $\bar{y}_e$ is a proportion in $(0, 1)$, a Beta likelihood is the natural choice:

$$\phi \sim \text{Gamma}(5, 0.1) \quad (\text{concentration/precision}) \quad (10)$$

$$\bar{y}_e \sim \text{Beta}(\mu_e \phi, (1 - \mu_e) \phi) \quad (11)$$

Beta Distribution Mean and Variance

If $Y \sim \text{Beta}(a, b)$, then:

$$\mathbb{E}[Y] = \frac{a}{a+b}, \quad \text{Var}[Y] = \frac{\mathbb{E}[Y](1 - \mathbb{E}[Y])}{\phi + 1}$$

In our parameterisation: $a = \mu_e \phi$, $b = (1 - \mu_e) \phi$, so:

$$\mathbb{E}[\bar{y}_e] = \frac{\mu_e \phi}{\mu_e \phi + (1 - \mu_e) \phi} = \mu_e$$

The Beta likelihood says: “the observed 21CS rate $\bar{y}_e$ should be close to the predicted rate μ_e , with concentration controlled by ϕ .” Larger ϕ = tighter fit. The prior $\phi \sim \text{Gamma}(5, 0.1)$ has mean 50, encoding moderate concentration.

2.5 Full Generative Story

BLFM Generative Process (for all episodes simultaneously)

1. For each feature f : draw precision $\alpha_f \sim \text{Gamma}(1, 1)$
2. For each feature f : draw weight $w_f \sim \mathcal{N}(0, 1/\alpha_f)$
3. Draw global intercept $b \sim \mathcal{N}(0, 1)$
4. Draw cell variance $\sigma_{\text{cell}} \sim \text{Half-Normal}(1)$
5. For each episode e : draw episode effect $u_e \sim \mathcal{N}(0, \sigma_{\text{cell}}^2)$
6. For each episode e : compute $\mu_e = \sigma(\mathbf{w}^\top \bar{\mathbf{x}}_e + u_e + b)$
7. Draw concentration $\phi \sim \text{Gamma}(5, 0.1)$
8. For each episode e : **observe** $\bar{y}_e \sim \text{Beta}(\mu_e \phi, (1 - \mu_e) \phi)$

2.6 Worked Numerical Example

Suppose MCMC returns these posterior means: $w_{\text{Beh}} = -1.5$, $w_{\text{Cog}} = -0.4$, $w_{\text{Con}} = +1.8$, $w_{\text{App}} = +2.1$ (all other $w_f \approx 0$), $b = -0.5$.

Episode A (mostly Behaviourist, low cognitive demand):

$$\bar{x}_{\text{Beh}} = 0.80, \quad \bar{x}_{\text{App}} = 0.05, \quad u_A = +0.1$$

$$\text{logit}(\mu_A) = (-1.5)(0.80) + (2.1)(0.05) - 0.5 + 0.1 = -2.0 \Rightarrow \hat{\mu}_A = \sigma(-2.0) = \mathbf{0.12}$$

Episode B (more Constructivist, higher Applying):

$$\bar{x}_{\text{Con}} = 0.30, \quad \bar{x}_{\text{App}} = 0.25, \quad u_B = -0.1$$

$$\text{logit}(\mu_B) = (1.8)(0.30) + (2.1)(0.25) - 0.5 - 0.1 = +0.465 \Rightarrow \hat{\mu}_B = \sigma(0.465) = \mathbf{0.61}$$

The model correctly assigns a much higher predicted 21CS rate to Episode B.

2.7 Code

```

1 def blfm_model(x_bar, y_bar, cell_idx, n_cells, P):
2     """
3     x_bar      : shape (E, 13)  -- mean feature vector per
                        episode
4     y_bar      : shape (E,)      -- observed 21CS rate per episode
5     cell_idx   : shape (E,)      -- Subject x Grade cell index
6     n_cells    : int             -- number of unique cells
7     P          : int             -- number of features (13)
8     """
9     # Step 1: ARD priors
10    # One precision alpha_f per feature
11    alpha = numpyro.sample("alpha",
12                           dist.Gamma(1.0, 1.0).expand([P]))
13    # Weight: w_f ~ N(0, 1/alpha_f)
14    # std = 1/sqrt(alpha) because Var(w_f) = 1/alpha_f
15    w = numpyro.sample("w",
16                       dist.Normal(jnp.zeros(P),
17                                   1.0 / jnp.sqrt(alpha)))
18    # Step 2: Global intercept
19    b = numpyro.sample("b", dist.Normal(0.0, 1.0))
20
21    # Step 3: Subject x Grade random effects
22    # sigma_cell controls how much episodes within a cell vary
23    sigma_cell = numpyro.sample("sigma_cell",
24                                dist.HalfNormal(1.0))
25    # u_cell[c] ~ N(0, sigma_cell^2) for each cell c
26    u_cell = numpyro.sample("u_cell",
27                            dist.Normal(jnp.zeros(n_cells),
28                                        sigma_cell * jnp.ones(n_cells)))
29    # Map each episode to its cell's random effect
30    u_e = u_cell[cell_idx]          # shape: (E,)
31
32    # Step 4: Predicted episode 21CS rate
33    # logit(mu_e) = w^T x_bar_e + u_e + b
34    logit_mu = jnp.dot(x_bar, w) + u_e + b
35    mu = jax.nn.sigmoid(logit_mu)   # mu_e in (0, 1)
36
37    # Step 5: Beta precision
38    # Gamma(5, 0.1): mean=50, represents moderate concentration
39    phi = numpyro.sample("phi", dist.Gamma(5.0, 0.1))

```

```

40
41     # Step 6: Beta likelihood
42     a = mu * phi + 1e-6          # alpha shape param
43     b_ = (1.0 - mu) * phi + 1e-6 # beta shape param
44     # Clamp y_bar to open interval: Beta requires (0,1)
45     # strictly
46     y_bar_clamped = jnp.clip(y_bar, 1e-4, 1.0 - 1e-4)
47     numpyro.sample("y_bar", dist.Beta(a, b_), obs=y_bar_clamped
48 )

```

Listing 1: BLFM: Full model definition in NumPyro

2.7.1 Key Implementation Details Explained

`1.0/jnp.sqrt(alpha)`

The standard deviation of w_f is $\alpha_f^{-1/2}$, because $\text{Var}(w_f) = 1/\alpha_f$. Always distinguish precision (α) from standard deviation.

`u_cell[cell_idx]`

`cell_idx[i]` stores which Subject×Grade cell episode i belongs to. This single indexing operation maps each of the E episodes to its cell’s random effect.

`jnp.clip(y_bar, 1e-4, 1-1e-4)`

The Beta distribution requires strict open-interval support. Any episode with $\bar{y}_e = 0$ (no 21CS at all) or $\bar{y}_e = 1$ would produce $\log \text{Beta}(-\infty)$. The clamp prevents this.

`1e-6 stabiliser on a and b_`

When $\mu \approx 0$ or $\mu \approx 1$, one of the Beta shape parameters would be nearly zero, causing numerical underflow. The small additive constant stabilises the computation.

2.8 Running MCMC

```

1 kernel_blfm = NUTS(blfm_model, target_accept_prob=0.90)
2 mcmc_blfm   = MCMC(kernel_blfm,
3                   num_warmup=500,      # adapt step size for 500
4                                     steps
5                   num_samples=1000,   # draw 1000 posterior
6                                     samples
7                   num_chains=2,       # run 2 independent
8                                     chains
9                   progress_bar=True)
10 mcmc_blfm.run(jax.random.PRNGKey(42),
11              x_bar_jax, y_bar_jax,
12              cell_idx_jax, n_cells, P)

```

Listing 2: BLFM: NUTS sampler configuration

NUTS (No-U-Turn Sampler) automatically tunes step size and trajectory length during warmup, making it far more efficient than basic Metropolis-Hastings for high-dimensional

posteriors. `target_accept_prob=0.90` means: aim for 90% acceptance, which NUTS achieves by tuning step sizes during warmup.

2.9 What BLFM Cannot Answer

BLFM Blind Spots

- What is the probability the classroom is productive *right now*, at line 47?
- Did the Constructivist move at line 32 *cause* a regime shift?
- How does the productive state evolve *across* the lesson arc?

These require a sequential model with a latent state. Enter PST.

3 Model 2: PST — Probabilistic Skill Tracing (Sequential BKT)

3.1 Core Question

PST Answers

How often does the classroom switch between productive and unproductive discourse? And what is the probability of being in the productive regime at each line, without knowing anything about the pedagogical content?

3.2 Hidden Markov Model (HMM) Primer

PST models the classroom as a two-state Hidden Markov Model. We have:

- A **latent state** $M_t \in \{0, 1\}$ at each line t — never directly observed.
- A **visible observation** $O_t \in \{0, 1\}$ — whether 21CS is active.

The HMM diagram below shows the structure:

3.3 PST Parameters

Symbol	Meaning	Prior
π_0	$P(M_1 = 1)$: probability episode starts in the productive state	Beta(2, 3)
p_L	Learning rate: $P(M_t=1 \mid M_{t-1}=0)$ per line	Beta(1.5, 20)
p_F	Forgetting rate: $P(M_t=0 \mid M_{t-1}=1)$ per line	Beta(1, 50)
p_S	Slip: $P(O_t=0 \mid M_t=1)$	Fixed = 0.015
p_G	Guess: $P(O_t=1 \mid M_t=0)$	Fixed = 0.009




Figure 1: PST graphical model. The discourse state M_t transitions between productive ($M=1$) and unproductive ($M=0$) with fixed probabilities p_L (learning) and p_F (forgetting), independent of pedagogical content. The observed 21CS activation O_t is emitted with slip probability p_S and guess probability p_G . No covariate vector $\mathbf{z}_t$ is present: content-blindness is intentional, isolating pure regime dynamics from pedagogical content. Compare IO-BKT (Fig. 2), which adds an orange $\mathbf{z}_t$ row and content-driven transition probabilities.

3.3.1 The Transition Matrix

At every time step, the probability of moving between states is constant (PST ignores content):

$$\mathbf{A} = \begin{pmatrix} P(M_t=0 \mid M_{t-1}=0) & P(M_t=1 \mid M_{t-1}=0) \\ P(M_t=0 \mid M_{t-1}=1) & P(M_t=1 \mid M_{t-1}=1) \end{pmatrix} = \begin{pmatrix} 1 - p_L & p_L \\ p_F & 1 - p_F \end{pmatrix} \quad (12)$$

The posterior mean $\hat{p}_L = 0.073$ means each line has only a 7.3% chance of transitioning from unproductive to productive. Regimes are sticky.

3.3.2 Emission Probabilities

$$P(O_t = 1 \mid M_t = 1) = 1 - p_S = 0.985 \quad (\text{productive} \Rightarrow \text{21CS almost always observed}) \quad (13)$$

$$P(O_t = 0 \mid M_t = 1) = p_S = 0.015 \quad (\text{"slip": productive but no 21CS}) \quad (14)$$

$$P(O_t = 1 \mid M_t = 0) = p_G = 0.009 \quad (\text{"guess": unproductive but 21CS fires}) \quad (15)$$

$$P(O_t = 0 \mid M_t = 0) = 1 - p_G = 0.991 \quad (\text{unproductive} \Rightarrow \text{no 21CS almost always}) \quad (16)$$

Why Call Them “Slip” and “Guess”?

These terms originate from BKT in tutoring systems (Corbett & Anderson, 1994):

- **Slip** ($p_S = 0.015$): Even when the classroom is genuinely in the productive regime, expert raters occasionally miss 21CS activation (momentary pause, annotation noise). This is rare: 21CS is almost always visible when the regime is truly productive.
- **Guess** ($p_G = 0.009$): Even in an unproductive regime, a single spontaneous student question might briefly activate 21CS. Also very rare.

Low slip and guess values mean the states are *nearly fully separable* by their emissions.

3.3.3 Why Are p_S and p_G Fixed?

This is a critical design decision. If we estimated p_S and p_G freely, we would face **label switching**: mathematically, swapping the labels “productive” and “unproductive” (and simultaneously swapping emission parameters) gives the exact same likelihood. MCMC would explore both interpretations, producing posterior means of $p_S \approx p_G \approx 0.5$ — meaningless.

Label Switching Visualised

Imagine the model has two possible solutions:

	Solution A (correct)	Solution B (flipped)
State 1 meaning	Productive	Unproductive
p_S (State 1 slip)	0.015	0.991
p_G (State 0 guess)	0.009	0.985
Likelihood	$\mathcal{L}$	$\mathcal{L}$ (identical!)

Both solutions have the same likelihood. Without content features to distinguish states, MCMC mixes freely between them. Fixing p_S and p_G at values from IO-BKT (which can identify them via content features) resolves the ambiguity.

3.4 The Forward Algorithm: Step-by-Step Derivation

The forward algorithm is the engine of PST. It computes the marginal likelihood $P(O_1, \dots, O_{T_e} \mid \theta)$ *exactly* in $O(T_e)$ time, by recursively integrating out the hidden states M_t .

3.4.1 Why Not Enumerate All State Sequences?

A brute-force approach would sum over all 2^{T_e} state sequences. For $T_e = 384$, this is $\sim 10^{115}$ terms — impossible. The forward algorithm avoids this by exploiting the **Markov property**: M_t depends only on M_{t-1} , not the full history.

3.4.2 Notation

Define the **filtered joint probability** at time t :

$$\alpha_t^{(m)} = P(M_t = m, O_1, O_2, \dots, O_t \mid \theta) \quad (17)$$

This is the probability that: (a) we have seen observations $O_1, \dots, O_t$, *and* (b) the state at time t is m .

3.4.3 Initialisation ($t = 1$)

$$\alpha_1^{(0)} = P(M_1 = 0) \cdot P(O_1 \mid M_1 = 0) = (1 - \pi_0) \cdot e_0(O_1) \quad (18)$$

$$\alpha_1^{(1)} = P(M_1 = 1) \cdot P(O_1 \mid M_1 = 1) = \pi_0 \cdot e_1(O_1) \quad (19)$$

where the emission shorthand is:

$$e_m(O_t) = \begin{cases} 1 - p_G & \text{if } m = 0, O_t = 0 \\ p_G & \text{if } m = 0, O_t = 1 \\ p_S & \text{if } m = 1, O_t = 0 \\ 1 - p_S & \text{if } m = 1, O_t = 1 \end{cases} \quad (20)$$

Normalise: $c_1 = \alpha_1^{(0)} + \alpha_1^{(1)} + \varepsilon$; then $\alpha_1^{(m)} \leftarrow \alpha_1^{(m)} / c_1$.

3.4.4 Recursion ($t \geq 2$): Predict–Update

Predict step — marginalise over the previous state:

$$\tilde{\alpha}_t^{(0)} = \alpha_{t-1}^{(0)} (1 - p_L) + \alpha_{t-1}^{(1)} p_F \quad (21)$$

$$\tilde{\alpha}_t^{(1)} = \alpha_{t-1}^{(0)} p_L + \alpha_{t-1}^{(1)} (1 - p_F) \quad (22)$$

Update step — weight by the emission probability:

$$\alpha_t^{(0)} \propto \tilde{\alpha}_t^{(0)} \cdot e_0(O_t) \quad (23)$$

$$\alpha_t^{(1)} \propto \tilde{\alpha}_t^{(1)} \cdot e_1(O_t) \quad (24)$$

Normalise with $c_t = \alpha_t^{(0)} + \alpha_t^{(1)} + \varepsilon$.

3.4.5 Log-Likelihood

$$\log P(O_1, \dots, O_{T_e} \mid \theta) = \sum_{t=1}^{T_e} \log c_t \quad (25)$$

This is the quantity NUTS maximises (in expectation) during MCMC.

3.5 Worked Numerical Forward Pass

Parameters: $\pi_0 = 0.4$, $p_L = 0.07$, $p_F = 0.02$, $p_S = 0.015$, $p_G = 0.009$.

Observations: $[O_1, O_2, O_3, O_4, O_5] = [1, 1, 0, 1, 1]$.

Emission probabilities: $e_0(O = 1) = 0.009$, $e_0(O = 0) = 0.991$, $e_1(O = 1) = 0.985$, $e_1(O = 0) = 0.015$.

t	O_t	$\tilde{\alpha}^{(0)}$	$\tilde{\alpha}^{(1)}$	$\alpha^{(0)}$ (normalised)	$\alpha^{(1)}$ (normalised)	c_t
1	1	—	—	0.0135	0.9865	0.3994
2	1	0.0323	0.9677	0.0003	0.9997	0.9535
3	0	0.0203	0.9797	0.5777	0.4223	0.0348
4	1	0.5457	0.4543	0.0109	0.9891	0.4524
5	1	0.0299	0.9701	0.0003	0.9997	0.9559

Reading the Forward Table

- $t = 1$: We observe $O_1 = 1$. Since $e_1(1) = 0.985 \gg e_0(1) = 0.009$, the 21CS observation strongly implies $M_1 = 1$: $\hat{P}(M_1 = 1) = 0.9865$.
- $t = 2$: Observing $O_2 = 1$ again. The predicted state before updating: $\tilde{\alpha}^{(0)} = (0.0135)(0.93) + (0.9865)(0.02) = 0.0323$. After the update with $O = 1$: $P(M_2 = 1)$ rises even higher to 0.9997.
- $t = 3$: $O_3 = 0$! No 21CS observed. The update weight for $M = 1$ is now $e_1(0) = p_S = 0.015$ (tiny), so $P(M_3 = 1)$ collapses to 0.4223. The classroom could have slipped out of the productive regime.
- $t = 4$: $O_4 = 1$ again. The state immediately recovers to $P(M_4 = 1) = 0.9891$. This is the near-permanent regime in action: one negative observation cannot permanently kill a productive state when p_S is small.

3.6 Prior Choices Explained

Prior	Mean	Reasoning
$\pi_0 \sim \text{Beta}(2, 3)$	0.40	Slight prior that episodes start unproductively. Most Indian classroom episodes (Behaviourist-dominant) begin with routine instruction.
$p_L \sim \text{Beta}(1.5, 20)$	0.07	Small learning rate expected. Consistent with near-permanent regime hypothesis: productive states are rare achievements, not the default.
$p_F \sim \text{Beta}(1, 50)$	0.02	Forgetting even rarer than learning. Once a productive regime is established, it tends to persist.

3.7 Key Result: Regime Persistence

$$\hat{p}_L = 0.073 \text{ } [0.067, 0.080], \quad \hat{p}_F = 0.021 \text{ } [0.012, 0.033] \quad (26)$$

Average lines per transition: $1/\hat{p}_L \approx 14$. A 384-line episode sees roughly $384 \times 0.073 \approx 28$ transitions total — but most are micro-fluctuations, not sustained regime changes.

3.8 Code

```

1 EPS      = 1e-7
2 PS_FIXED = 0.015    # from IO-BKT posterior (fixes label
   switching)
3 PG_FIXED = 0.009    # from IO-BKT posterior
4
5 def pst_forward_single(obs_e, mask_e, pi0, p_L, p_F):
6     """

```

```

7      obs_e : shape (Tmax,) -- 21CS observations (0/1), padded
8      mask_e : shape (Tmax,) -- 1 for real lines, 0 for padding
9      pi0    : scalar      -- P(M_1=1)
10     p_L    : scalar      -- learning rate (0->1 per line)
11     p_F    : scalar      -- forgetting rate (1->0 per line)
12     """
13     pS = PS_FIXED
14     pG = PG_FIXED
15
16     # Initialisation at t=0
17     o0 = obs_e[0]
18     # Emission:  $P(O_1 | M_1=0)$  and  $P(O_1 | M_1=1)$ 
19     e0_0 = jnp.where(o0 == 1., pG, 1. - pG)
20     e1_0 = jnp.where(o0 == 1., 1. - pS, pS)
21     #  $\alpha_1^{(0)}$  and  $\alpha_1^{(1)}$ , unnormalised
22     a0 = (1. - pi0) * e0_0
23     a1 = pi0 * e1_0
24     # Normalise and log the normalisation constant
25     c0 = a0 + a1 + EPS
26     a0 = a0 / c0; a1 = a1 / c0
27     ll0 = jnp.log(c0) * mask_e[0] # 0 if padding line
28
29     # Recursion: t=1, 2, ..., Tmax-1
30     def step(carry, t):
31         a0, a1, ll = carry
32         o_t = obs_e[t]
33         m_t = mask_e[t] # padding guard
34
35         # PREDICT: marginalise over  $M_{t-1}$ 
36         pred0 = a0 * (1. - p_L) + a1 * p_F # Eq.(12)
37         pred1 = a0 * p_L + a1 * (1. - p_F) # Eq.(13)
38
39         # Emission probabilities for current observation
40         e0 = jnp.where(o_t == 1., pG, 1. - pG)
41         e1 = jnp.where(o_t == 1., 1. - pS, pS)
42
43         # UPDATE: weight by emission
44         b0 = pred0 * e0
45         b1 = pred1 * e1
46         c = b0 + b1 + EPS
47
48         # Accumulate log-likelihood (zero for padding)
49         return (b0/c, b1/c, ll + jnp.log(c) * m_t), None
50
51     # lax.scan: JIT-compiled loop, runs efficiently on GPU
52     (_, _, ll_final), _ = lax.scan(
53         step, (a0, a1, ll0), jnp.arange(1, Tmax)
54     )
55     return ll_final

```

Listing 3: PST: Forward algorithm for one episode

```

1  # Vectorise forward pass: apply to all E episodes at once
2  pst_forward_all = vmap(pst_forward_single,
3                          in_axes=(0, 0, None, None, None))
4
5  def pst_model(obs, mask):
6      """
7      Only 3 parameters are estimated -- content-free by design.
8      """
9      # Beta(2,3): mean=0.40, slight prior for unproductive start
10     pi0 = numpyro.sample("pi0", dist.Beta(2.0, 3.0))
11     # Beta(1.5,20): mean=0.07, small learning rate expected
12     p_L = numpyro.sample("p_L", dist.Beta(1.5, 20.0))
13     # Beta(1,50): mean=0.02, forgetting even rarer
14     p_F = numpyro.sample("p_F", dist.Beta(1.0, 50.0))
15
16     # Forward pass on all episodes simultaneously
17     ll = pst_forward_all(obs, mask, pi0, p_L, p_F)
18     # numpyro.factor adds log-prob directly to the model
19     numpyro.factor("seq_loglik", jnp.sum(ll))
20
21     # Run NUTS MCMC
22     kernel_pst = NUTS(pst_model, target_accept_prob=0.90)
23     mcmc_pst = MCMC(kernel_pst, num_warmup=500,
24                     num_samples=1000, num_chains=2)
25     mcmc_pst.run(jax.random.PRNGKey(42), obs_jax, mask_jax)

```

Listing 4: PST: Model and MCMC

3.8.1 Two Technical Details

jax.lax.scan A standard Python for loop would unroll to $T_{\max}$ Python-level operations, each launching a separate GPU kernel — very slow. `lax.scan` compiles the entire loop into a single XLA kernel, running the forward pass on GPU in milliseconds.

The mask trick All episodes are zero-padded to a uniform length $T_{\max}$. The mask ensures padding lines contribute $\log(c_t) \times 0 = 0$ to the log-likelihood. Without masking, the algorithm would treat padding zeros as real “no-21CS” observations, artificially inflating the estimated p_F .

3.9 Why PST’s AUC is Misleading

PST achieves $\text{AUC} = 0.928$, higher than IO-BKT’s 0.873. But this is not a genuine win.

PST AUC is Autocorrelation, Not Insight

PST’s one-step-ahead prediction of O_t is:

$$\hat{P}(O_t = 1 \mid O_1, \dots, O_{t-1}) \approx O_{t-1}$$

The correlation $\text{corr}(P_t, O_{t-1}) = 0.998$ confirms this. PST is essentially saying “the next observation looks like the current one,” which is a tautology when regimes barely change. There is no pedagogical insight: IO-BKT’s 0.873, achieved with *content features alone* (not the observation history), is the honest comparison.

4 Model 3: IO-BKT — Input-Output Bayesian Knowledge Tracing

4.1 Core Question

IO-BKT Answers

Which EP/KD/CP codes are most strongly associated with transitions into and out of the productive discourse regime at each utterance? The transition coefficients β_{on} and β_{off} quantify these associations from observational data; causal interpretation requires replication with experimental or quasi-experimental designs.

4.2 The Key Innovation: Content-Driven Transitions




Figure 2: IO-BKT graphical model. At each line t , the observed EP/KD/CP covariate vector $\mathbf{z}_t$ modulates both the 0→1 transition probability (ignition, via β_{on}) and the 1→0 transition probability (dropout, via β_{off}), making transitions content-driven rather than fixed. The latent productive-discourse state M_t is inferred via the forward algorithm; 21CS activation O_t is the observed emission. Compare PST (Fig. 1), which has no $\mathbf{z}_t$ row and uses constant transition probabilities.

PST uses a constant transition probability: every line has the same p_L (chance of $0 \rightarrow 1$)

regardless of whether the teacher is asking “What is the formula?” or “Now apply this to design a bridge.” This is clearly wrong.

IO-BKT replaces constant transitions with **sigmoid-linear functions of $\mathbf{z}_t$** :

$$P(M_t = 1 \mid M_{t-1} = 0, \mathbf{z}_t) = \sigma(\boldsymbol{\beta}_{\text{on}}^\top \mathbf{z}_t + u_e) \quad (27)$$

$$P(M_t = 0 \mid M_{t-1} = 1, \mathbf{z}_t) = \sigma(\boldsymbol{\beta}_{\text{off}}^\top \mathbf{z}_t + u_e) \quad (28)$$

The Two Coefficient Vectors Have Distinct Roles

- **$\boldsymbol{\beta}_{\text{on}}$ (ignition coefficients):** $\beta_{\text{on},f} > 0$ means feature f *promotes* entry into the productive state from unproductive. $\beta_{\text{on},f} < 0$ means feature f *suppresses* ignition.
- **$\boldsymbol{\beta}_{\text{off}}$ (dropout coefficients):** $\beta_{\text{off},f} > 0$ means feature f *promotes* dropout from productive to unproductive. $\beta_{\text{off},f} < 0$ means feature f *prevents* dropout (sustains the productive state).

A feature can have *opposite* signs in $\boldsymbol{\beta}_{\text{on}}$ and $\boldsymbol{\beta}_{\text{off}}$ — for example, Applying has $\beta_{\text{on}} = +3.41$ (ignites) and $\beta_{\text{off}} = -3.40$ (sustains). A simple regression or correlation cannot separate these roles.

4.3 Full Model Specification

4.3.1 Priors

$$\boldsymbol{\beta}_{\text{on}}, \boldsymbol{\beta}_{\text{off}} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_{14}) \quad (29)$$

$$\pi_0 \sim \text{Beta}(2, 3) \quad (30)$$

$$p_S \sim \text{Beta}(2, 30) \text{ [mean} \approx 0.06] \quad (31)$$

$$p_G \sim \text{Beta}(1.5, 30) \text{ [mean} \approx 0.05] \quad (32)$$

$$\sigma_{\text{cell}} \sim \text{Half-Normal}(1) \quad (33)$$

$$u_e \sim \mathcal{N}(0, \sigma_{\text{cell}}^2) \quad (34)$$

Unlike PST, IO-BKT *estimates* p_S and p_G . The content features break label-switching symmetry: by the time MCMC runs, it knows Applying (+3.41) is associated with ignition, so the “productive” label attaches to the state where high-CP lines drive entry — there is no ambiguity.

4.3.2 Emission Model

Unchanged from PST (Eqs. 13-16), but now p_S and p_G are *estimated*:

$$P(O_t = 1 \mid M_t = 1) = 1 - p_S, \quad P(O_t = 1 \mid M_t = 0) = p_G \quad (35)$$

The posterior means are $\hat{p}_S = 0.015$, $\hat{p}_G = 0.009$ — confirming PST’s fixed values were appropriate.

4.3.3 Forward Algorithm with Time-Varying Transitions

The forward algorithm is identical in structure to PST, but now p_L and p_F are replaced by line-specific values $p_{\text{on},t}$ and $p_{\text{off},t}$:

Predict step:

$$\tilde{\alpha}_t^{(0)} = \alpha_{t-1}^{(0)} (1 - p_{\text{on},t}) + \alpha_{t-1}^{(1)} p_{\text{off},t} \quad (36)$$

$$\tilde{\alpha}_t^{(1)} = \alpha_{t-1}^{(0)} p_{\text{on},t} + \alpha_{t-1}^{(1)} (1 - p_{\text{off},t}) \quad (37)$$

Update step: same as PST (weight by emission). The key difference: every line has its own $p_{\text{on},t}$ and $p_{\text{off},t}$, computed from $\mathbf{z}_t$ via sigmoid.

4.4 Worked Numerical Example: Transition Probabilities

Using the paper's posterior means (β_{on} and β_{off} from Table II, intercept $u_e = 0$):

Line A — Behaviourism + Remembering: $\mathbf{z}_A = [\underbrace{1, 0, 0, 0, 0, 0, 0}_{\text{EP}}, \underbrace{1, 0, 0, 0, 0, 0}_{\text{KD}}, \underbrace{1, 0, 0, 0, 0, 0}_{\text{CP}}]$

$$\beta_{\text{on}}^\top \mathbf{z}_A = (-1.50)_{\text{Beh}} + (-1.15)_{\text{Rem}} + (-1.78)_{\text{int}} = -4.43$$

$$p_{\text{on},A} = \sigma(-4.43) = \mathbf{0.012} \text{ (1.2\% ignition chance)}$$

Line B — Constructivism + Applying: $\mathbf{z}_B = [\underbrace{0, 0, 1, 0, 0, 0, 0}_{\text{EP}}, \underbrace{0, 0, 0, 0, 0, 0}_{\text{KD}}, \underbrace{0, 0, 1, 0, 0, 0}_{\text{CP}}]$

$$\beta_{\text{on}}^\top \mathbf{z}_B = (+1.25)_{\text{Con}} + (+3.41)_{\text{App}} + (-1.78)_{\text{int}} = +2.88$$

$$p_{\text{on},B} = \sigma(+2.88) = \mathbf{0.947} \text{ (94.7\% ignition chance)}$$

What This Means for a Teacher

If a teacher is giving a Behaviourist recall drill (Line A type), each line has only a 1.2% chance of igniting a productive discourse regime. But if the teacher shifts to a Constructivist application task (Line B type), the probability jumps to 94.7%. The IO-BKT coefficients operationalise precisely *which* moves matter and by how much.

4.5 Interpreting Table II: The Posterior Coefficients

Ignition (β_{on})			Dropout prevention / promotion (β_{off})		
Feature	Coeff.	$ \mu > 2\text{SD}$	Feature	Coeff.	$ \mu > 2\text{SD}$
Applying	+3.41	★	Applying	-3.40	★
Constructivism	+1.25		Constructivism	-2.47	★
Behaviourism	-1.50	★	Cognitivism	-1.31	★
Remembering	-1.15	★	Remembering	+1.23	★
Procedural	-0.89	★	Understanding	-0.80	★
Intercept	-1.78	★	Intercept	+0.44	

Reading each coefficient:**Applying**, $\beta_{\text{on}} = +3.41^*$ and $\beta_{\text{off}} = -3.40^*$

The dominant lever in both directions. Applying-level cognitive demand (Bloom’s level 3: “use knowledge in a new situation”) both *ignites* the productive state and *sustains* it once established. It is the most impactful feature in the model.

Constructivism, $\beta_{\text{off}} = -2.47^*$

Once the classroom is productive, Constructivist moves (Vygotsky zone-of-proximal-development scaffolding, collaborative discovery) are the strongest *prevention of dropout*. Note: $\beta_{\text{on}} = +1.25$ for Constructivism is *not* credibly non-zero ($< 2\text{SD}$), meaning Constructivism sustains but does not reliably *ignite* (likely due to its extreme rarity in the corpus: $< 5\%$ of lines).

Behaviourism, $\beta_{\text{on}} = -1.50^*$

Operant-conditioning-style instruction (drill, rote recall, reward/punish) actively *suppresses* entry into the productive regime.

Remembering, $\beta_{\text{on}} = -1.15^*$ and $\beta_{\text{off}} = +1.23^*$

Pure recall tasks both suppress ignition *and* are the most reliable dropout trigger. If the classroom is productive and the teacher reverts to “remember this definition,” the productive state is at serious risk.

Intercept, $\beta_{\text{on}} = -1.78^*$

The baseline ignition probability (when all features are zero) is $\sigma(-1.78) = 0.144$. **Classrooms default toward unproductive discourse:** productivity requires active effort.

4.6 Code

```

1 def forward_single(obs_e, mask_e, pon_e, poff_e, pi0, pS, pG):
2     """
3     pon_e   : shape (Tmax,) -- P(M_t=1|M_{t-1}=0, z_t), per line
4     poff_e  : shape (Tmax,) -- P(M_t=0|M_{t-1}=1, z_t), per line
5     These vary each line -- the key difference from PST!
6     """
7     # Initialise
8     o0 = obs_e[0]
9     e0_0 = jnp.where(o0==1., pG, 1.-pG)      # P(O1|M1=0)
10    e1_0 = jnp.where(o0==1., 1.-pS, pS)      # P(O1|M1=1)
11    a0 = (1.-pi0) * e0_0
12    a1 = pi0 * e1_0
13    c0 = a0 + a1 + EPS
14    a0 /= c0; a1 /= c0
15    ll0 = jnp.log(c0) * mask_e[0]
16
17    def step(carry, t):
18        a0, a1, ll = carry
19        o_t = obs_e[t]
20        m_t = mask_e[t]

```



```

22         sigma_cell * jnp.ones(n_cells)))
23
24     # HMM parameters
25     pi0 = numpyro.sample("pi0", dist.Beta(2.0, 3.0))
26     pS = numpyro.sample("pS", dist.Beta(2.0, 30.0)) # mean
        ~0.06
27     pG = numpyro.sample("pG", dist.Beta(1.5, 30.0)) # mean
        ~0.05
28
29     # Compute time-varying transition probabilities
30     u_e = u_cell[cell_idx] # shape: (E,)
31     # eta_on[e,t] = sum_p (Z[e,t,p] * b_on[p]) + u_e[e]
32     # "etp,p->et" means: sum over last axis (features p)
33     eta_on = jnp.einsum("etp,p->et", Z, b_on) + u_e[:, None]
34     eta_off = jnp.einsum("etp,p->et", Z, b_off) + u_e[:, None]
35     # Squeeze through sigmoid to get probabilities
36     p_on = jax.nn.sigmoid(eta_on) # shape: (E, Tmax)
37     p_off = jax.nn.sigmoid(eta_off) # shape: (E, Tmax)
38     # p_on[e,t] = P(M_t=1 | M_{t-1}=0, z_{e,t})
39     # p_off[e,t] = P(M_t=0 | M_{t-1}=1, z_{e,t})
40
41     # Run forward algorithm over all episodes
42     ll_per_episode = forward_all(obs, mask, p_on, p_off,
43                                 pi0, pS, pG)
44     numpyro.factor("seq_loglik", jnp.sum(ll_per_episode))

```

Listing 6: IO-BKT: NumPyro model with hierarchical effects

4.6.1 The einsum Operation Decoded

```

1 eta_on = jnp.einsum("etp,p->et", Z, b_on) + u_e[:, None]

```

This computes for all episodes e and time steps t simultaneously:

$$\eta_{\text{on},e,t} = \sum_{p=1}^P Z_{e,t,p} \cdot \beta_{\text{on},p} + u_e = \beta_{\text{on}}^\top \mathbf{z}_{e,t} + u_e$$

The index string "etp,p->et" means: *sum over dimension p (features), keeping dimensions e (episodes) and t (time steps)*. This replaces a nested Python loop with a single GPU matrix operation. Without `einsum`, you would write:

```

1 # Equivalent but slow:
2 for e in range(E):
3     for t in range(Tmax):
4         eta_on[e,t] = sum(Z[e,t,:] * b_on) + u_e[e]

```

4.6.2 MCMC Configuration

```

1 kernel = NUTS(

```

```

2     iobkt_model,
3     target_accept_prob=0.95,      # higher than BLFM's 0.90
4     max_tree_depth=10,           # allow longer HMC trajectories
5 )
6 mcmc = MCMC(
7     kernel,
8     num_warmup=500,
9     num_samples=1000,
10    num_chains=2,
11    chain_method='sequential', # avoid GPU memory doubling
12    progress_bar=True,
13 )

```

Listing 7: IO-BKT: NUTS with tighter tolerances for complex geometry

Why target_accept_prob=0.95? IO-BKT’s posterior has a more complex, curved geometry than BLFM’s, because the 28 coefficients (β_{on} and β_{off}) interact with the forward algorithm non-linearly. A higher target acceptance forces NUTS to take smaller, more careful steps (via a smaller step size), reducing the number of divergent transitions at the cost of slower exploration. This is the standard recommendation when you see many divergences with the default 0.80 or 0.90.

4.7 Label Switching: Detection and Correction

```

1 pS_mean = float(np.array(samples["pS"]).mean())
2 if pS_mean > 0.5:
3     # The model learned the states in REVERSE order.
4     # "State 1" actually behaves like the unproductive state.
5     print("Label switching detected -- flipping states")
6
7     # Flip all coefficient signs: if b_on drove 0->1 transition
8     # in the wrong direction, negating it flips the roles.
9     samples["b_on"] = -samples["b_on"]
10    samples["b_off"] = -samples["b_off"]
11
12    # Correct emission parameters.
13    # Original "state 1" had slip-like pS = 1 - pG_correct.
14    # After flipping: new pS = 1 - old pG, new pG = 1 - old
15    # pS
16    pS_old = samples["pS"].copy()
17    pG_old = samples["pG"].copy()
18    samples["pS"] = 1.0 - pG_old
19    samples["pG"] = 1.0 - pS_old

```

Listing 8: Label switching check and correction after MCMC

Why does this happen? If MCMC’s random initialisation starts near the “wrong” mode (where state 0 is actually more productive than state 1), it may converge to the label-swapped solution. The diagnostic: $\hat{p}_S > 0.5$ would mean the “productive” state has a 50%+ chance of producing no 21CS — absurd. Negating the coefficients and swapping

emission parameters corrects this.

4.8 The Full-Length Normalisation (Tmax Extension)

The training used $T_{\max} = 600$ (92nd percentile). Some episodes (e.g. Episode 37 with 921 lines) were truncated. The normalisation extension re-runs *only the forward pass* with $T_{\max} = 950$, keeping all MCMC-estimated parameters unchanged.

```

1  # b_on_m, b_off_m etc. are the posterior MEANS (fixed, from
   MCMC)
2  # We just extend the input arrays and re-apply the forward pass
   .
3
4  for i in range(E):
5      T_new = int(mask_full[i].sum())    # now using all lines
6      u_e   = float(u_cell_m[cell_idx[i]])
7
8      # Time-varying transitions for ALL T_new lines
9      # (same formula as training, just more timesteps)
10     pon_e  = sigmoid(Z_full[i, :T_new] @ b_on_m + u_e)
11     poff_e = sigmoid(Z_full[i, :T_new] @ b_off_m + u_e)
12
13     # Run the same forward algorithm over extended sequence
14     probs = forward_pass(obs_full[i, :T_new], pon_e, poff_e,
        ...)
```

Listing 9: Forward pass extension without re-running MCMC

This is valid because the IO-BKT coefficients describe how EP/KD/CP features drive transitions *in general*, they are not specific to the first 600 lines. Applying them to lines 601-950 is legitimate extrapolation, not re-fitting.

5 Experimental Setup

Metrics. Following Scarlatos et al. [1], three metrics have been used namely Accuracy (Acc), AUC, and F1 at line level ($n=18,506$; first line of each episode excluded as no prior-observation context available). AUC is used as the primary metric. Predictions are one-step-ahead: $P(O_t=1 \mid O_{1:t-1})$ computed before O_t is revealed. Verified by lag-correlation test: $\text{corr}(P_t, O_{t-1})=0.998$, $\text{corr}(P_t, O_t)=0.766$ for PST with no leakage.

Evaluation: Bayesian models use posterior predictive distributions (regularising priors mitigate overfitting). DKT uses LOO-CV. BKT uses in-sample forward smoother.

6 Results and Discussions

6.1 Baselines

The three baselines span the design space and each serves a specific evidentiary role; they were not selected for convenience but to answer three distinct questions about what the proposed framework adds over existing tools.

Frequentist BKT (pyBKT) is the field’s default sequential KT method. We run BKT aggregate (one KC across all lines) and BKT per-KC (dominant EP as the KC label), each with 5 random initialisations. Its role is diagnostic: if standard BKT works, the problem is already solved. That both variants collapse to $\text{AUC} = 0.500$ is itself a result – it establishes that near-permanent regimes are precisely the failure mode classical KT is not designed for, because the filtered posterior assigns near-constant within-episode probabilities when p_L is small, losing all within-episode discrimination.

DKT-LSTM (content-only): 1-layer LSTM (hidden = 64, dropout = 0.2), input , Adam ($\text{lr} = 10^{-3}$, $\text{wd} = 10^{-4}$), 30 epochs. It uses identical inputs to IO-BKT but replaces the interpretable HMM with a recurrent black box. DKT is evaluated by leave-one-episode-out CV (LOO, 48 folds), giving it a *stricter* evaluation than the in-sample Bayesian models. This asymmetry is principled: Bayesian priors act as regularisation, making in-sample posterior predictive distributions analogous to cross-validated frequentist estimates. That IO-BKT wins by +0.201 AUC *despite* DKT’s evaluation advantage establishes that principled Bayesian structure outperforms neural flexibility at this corpus size. PSIS-LOO comparison across all models is not reported because PST and IO-BKT sequence log-likelihoods (~ 390 per-line contributions) are incommensurable with BLFM’s single Beta observation per episode.

BLFM (§2) is the exchangeable Bayesian baseline from our own framework. Its inclusion provides a within-framework ablation: the AUC gap to IO-BKT (+0.101) isolates precisely what temporal structure with content-driven transitions contributes over the episode-level sufficient statistic.

6.2 What Each AUC Gap Means

BLFM $\rightarrow$ IO-BKT: +0.101 Adding temporal structure with content-driven transitions is worth 10 AUC points. The episode-mean sufficient statistic throws away real sequential information.

Table 1: KT comparison. Metrics at line level ($n=18,506$), first line per episode excluded [1].
 $\dagger$ = uses observation history; $\ddagger$ = LOO-CV (48 folds).

Method	Type	Acc	AUC	F1
<i>Group A: Content-only (EP+KD+CP, no observation history)</i>				
IO-BKT (Bayesian)	Seq., in-sample	0.878	0.873	0.875
BLFM (Bayesian)	Exch., in-sample	0.700	0.772	0.669
DKT-LSTM	Seq., LOO $\ddagger$	0.623	0.672	0.575
<i>Group B: Uses observation history (AUC reflects autocorrelation§)</i>				
PST (Bayesian) $\dagger$	Seq., in-sample	0.880	0.928	0.876
BKT agg. (freq.) $\dagger$	Seq., in-sample	0.518	0.500	0.000
BKT per-KC (freq.) $\dagger$	Seq., in-sample	0.517	0.500	0.000

§ PST $\text{corr}(P_t, O_{t-1})=0.998$: AUC from regime persistence.
 BKT AUC=0.5: constant within-episode predictions (regime barely changes).

DKT $\rightarrow$ IO-BKT: +0.201 Despite IO-BKT being evaluated in-sample while DKT uses stricter LOO-CV, IO-BKT wins. Principled Bayesian priors serve as implicit regularisation, outperforming a data-hungry neural network on a small ($E = 48$ episodes), diverse corpus.

IO-BKT $\rightarrow$ PST: +0.055 *Not* a genuine improvement. This gap is pure autocorrelation. On data where regimes are near-permanent, predicting “same as last line” trivially achieves high AUC.

Group A (content-only, fair comparison to DKT). IO-BKT AUC = 0.873 surpasses BLFM by +0.101 and DKT by +0.201. The BLFM $\rightarrow$ IO-BKT gap isolates the contribution of temporal structure with content-driven transitions over the exchangeable assumption. IO-BKT beats DKT evaluated under the stricter LOO protocol despite IO-BKT being evaluated in-sample, underscoring the value of principled Bayesian structure on small, diverse data.

Group B (observation history). PST AUC = 0.928 reflects autocorrelation, not content: $\text{corr}(P_t, O_{t-1})=0.998$ confirms predictions come almost entirely from the previous observation. This is the measurement of regime persistence, not pedagogical insight. Frequentist BKT collapses to AUC = 0.500 because its filtered posterior assigns near-constant values within each episode, the regime barely changes, losing within-episode discrimination.

6.3 PST: Regime Persistence Finding


$$\begin{aligned}\hat{p}_L &= 0.073 \text{ [95\% CI : 0.067, 0.080]}, \\ \hat{p}_F &= 0.021 \text{ [95\% CI : 0.012, 0.033]}\end{aligned}\tag{38}$$

In this corpus, the discourse state transitions roughly once every 14 lines on average. IO-BKT productive-state trajectories (Fig. ??) for science and mathematics topics, confirm the finding visually: For science, the episodes that open in the productive regime maintain a mean $P(M_t=1) \approx 0.65$ throughout; unproductive episodes remain near 0.25 and rarely recover. Similarly, for mathematics, only 3 unproductive episodes. But crucially, they

fluctuate dramatically (0.0 to 0.8). Mathematics unproductive episodes show recovery potential that Science unproductive episodes do not.

6.4 The Three-Stage AI Coaching Pipeline

The paper’s models together define a deployable real-time coaching system:



Stage	Details
1. Annotate	Fine-tune an open-source LLM (e.g. Llama 3) on the 19,181 expert-annotated lines to predict EP/KD/CP. Reduces cost from ~ 3 person-hours per episode to seconds. Classroom audio via Whisper completes the pipeline.
2. Detect	Run PST’s forward algorithm on the first 30 lines of the episode to classify whether the lesson is on a productive trajectory.
3. Recommend	Use IO-BKT’s β_{on} and β_{off} as a portable decision model in a teacher-facing app: • Ignite: increase Applying-level demand (+3.41) • Sustain: deploy Constructivism (−2.47 on dropout) • Avoid: Remembering-level drops (+1.23 dropout trigger)

6.5 Summary: The Three Models Side by Side

Property	BLFM	PST	IO-BKT
Data unit	Episode mean $\bar{\mathbf{x}}_e$	Line sequence $O_{1:T}$	Lines $O_{1:T}$ + features $\mathbf{z}_{1:T}$
Latent state	None (regression)	$M_t \in \{0, 1\}$	$M_t \in \{0, 1\}$
Transition	N/A	Fixed p_L, p_F (content-free)	$\sigma(\boldsymbol{\beta}^\top \mathbf{z}_t)$ per line
Emission	Beta regression	Fixed p_S, p_G	Estimated p_S, p_G
Parameters	$\mathbf{w}, b, \phi, \sigma_{\text{cell}}$	π_0, p_L, p_F	$\boldsymbol{\beta}_{\text{on}}, \boldsymbol{\beta}_{\text{off}}, \pi_0, p_S, p_G, \sigma_{\text{cell}}$
Key question	Does content mix predict 21CS rate?	How stable are discourse regimes?	Which moves ignite/extinguish productive talk?
Key finding	21CS rate $\propto$ content mix	$\hat{p}_L = 0.073$: regimes near-permanent	Applying: +3.41; Constructivism: -2.47; Remembering: +1.23
Runtime	< 5 min (CPU)	~ 15 min (GPU)	~ 180 min (GPU)
AUC	0.772	0.928 [†]	0.873
Use when	Bulk analysis, no per-line prediction needed	Measuring regime stability, early classification	Identifying causal moves, AI coaching

[†] Inflated by autocorrelation ($\text{corr}(P_t, O_{t-1}) = 0.998$).

6.6 Conclusion

The three models form a deliberate progression:

1. **BLFM** establishes the baseline: yes, content predicts 21CS, and the AUC gap (0.772 vs 0.873) quantifies exactly what is lost by discarding temporal order.
2. **PST** measures the regime dynamics: transitions are rare ($\hat{p}_L = 0.073$, $\hat{p}_F = 0.021$). The classroom spends most of its time in one regime. This empirically *justifies* BLFM’s exchangeability assumption — and motivates early-detection tools (classify by line 30; the regime will likely hold).
3. **IO-BKT** provides the actionable layer: *what causes* transitions. Applying-level cognitive demand (+3.41) is the primary ignition driver. Constructivism (−2.47 on dropout) sustains productive discourse once established. Remembering (+1.23) is

the most reliable regime killer. These coefficients are not descriptive statistics — they are a portable, deployable decision model for the next generation of AI teaching coaches.

The near-permanent regime finding is specific to this predominantly Behaviourist, India-dominant corpus and requires replication in Socratic, inquiry-based, or dialogic classrooms before strong deployment claims are made. Subject to that caveat, these three models provide the first principled Bayesian framework for tracking classroom-level productive discourse — the missing upstream layer that student-level knowledge tracing requires.

6.7 Knowledge Change Curves

6.7.1 Motivation

The IO-BKT transition coefficients and (Table 4.5) capture the per-line effect of each feature on regime transitions. They cannot, however, reveal whether repeated use of a feature within a lesson produces a *cumulative* shift in the productive-discourse regime that differs from its instantaneous effect. Knowledge change curves address this by tracking the IO-BKT posterior productive-state probability across successive occurrences of each feature within an episode, adapting the knowledge-change curve methodology of Scarlatos et al. [1] from individual KC mastery in tutoring dialogues to classroom-level discourse regimes.

6.7.2 Notation

Let $e \in \{1, \dots, E\}$ index episodes and $t \in \{1, \dots, T_e\}$ index lines within episode e . The IO-BKT forward algorithm produces, for every line t in episode e , the filtered posterior:

$$P^{(e)}(M_t=1) = P\left(M_t = 1 \mid O_1^{(e)}, \dots, O_t^{(e)}, \mathbf{z}_1^{(e)}, \dots, \mathbf{z}_t^{(e)}, \hat{\theta}\right) \quad (39)$$

where $\hat{\theta} = \{\hat{\cdot}, \hat{\pi}_0, \hat{p}_S, \hat{p}_G, \hat{\sigma}_{\text{cell}}\}$ are the IO-BKT posterior means, and $O_t^{(e)} \in \{0, 1\}$ is the observed 21CS activation at line t in episode e . This is a *latent* quantity: M_t is never directly observed.

6.7.3 Episode Opening Baseline

The baseline for episode e is the filtered posterior at the *first line of the episode*:

$$P^{(e)}(M_1=1) = \frac{\hat{\pi}_0 \cdot P(O_1^{(e)} \mid M_1 = 1)}{\hat{\pi}_0 \cdot P(O_1^{(e)} \mid M_1 = 1) + (1 - \hat{\pi}_0) \cdot P(O_1^{(e)} \mid M_1 = 0)} \quad (40)$$

where $P(O_1 = 1 \mid M_1 = 1) = 1 - \hat{p}_S$ and $P(O_1 = 1 \mid M_1 = 0) = \hat{p}_G$. The baseline $P^{(e)}(M_1=1)$ is a model prediction, not a data reading: it reflects $\hat{\pi}_0$ immediately updated on the first observed 21CS activation $O_1^{(e)}$.

6.7.4 Feature Occurrence Indexing

For a binary feature $f \in \{1, \dots, 13\}$ (any EP, KD, or CP code), define the ordered set of lines in episode e where feature f is active:

$$\mathcal{T}_{e,f} = \{t : z_{t,f}^{(e)} = 1\} = \{t_{e,f}^{(1)}, t_{e,f}^{(2)}, \dots\} \quad (41)$$

in ascending order of t . The j th element $t_{e,f}^{(j)}$ is the line at which feature f appears for the j th time in episode e . Define $J_{e,f} = |\mathcal{T}_{e,f}|$ as the total number of occurrences of feature f in episode e .

6.7.5 Per-Occurrence Change from Opening

At the j th occurrence of feature f in episode e , the change in productive-state probability from the episode opening is:

$$\Delta P_{e,f}^{(j)} = P^{(e)}(M_{t_{e,f}^{(j)}}=1) - P^{(e)}(M_1=1) \quad (42)$$

$\Delta P_{e,f}^{(j)} > 0$ means the classroom productive-state probability is above its own opening level at the j th occurrence of feature f ; $\Delta P_{e,f}^{(j)} < 0$ means it is below. Both quantities on the right-hand side of Eq. (42) are IO-BKT forward-algorithm outputs: the curve lives entirely in latent state space.

6.7.6 The Knowledge Change Curve

For occurrence index j and feature f , let $\mathcal{E}_f^{(j)} = \{e : J_{e,f} \geq j\}$ be the set of episodes containing at least j occurrences of feature f . The knowledge change curve value at occurrence j is:

$$\boxed{\overline{\Delta P}_f^{(j)} = \frac{1}{|\mathcal{E}_f^{(j)}|} \sum_{e \in \mathcal{E}_f^{(j)}} \Delta P_{e,f}^{(j)}} \quad (43)$$

with standard error:

$$\text{SE}_f^{(j)} = \frac{1}{\sqrt{|\mathcal{E}_f^{(j)}|}} \sqrt{\frac{1}{|\mathcal{E}_f^{(j)}| - 1} \sum_{e \in \mathcal{E}_f^{(j)}} \left(\Delta P_{e,f}^{(j)} - \overline{\Delta P}_f^{(j)} \right)^2} \quad (44)$$

The curve is plotted for $j = 1, \dots, 15$. The label $n_{\text{eps}} \leq N$ in each panel of Fig. ?? reports $|\mathcal{E}_f^{(1)}|$, the number of contributing episodes at occurrence 1; this decreases with j as shorter episodes drop out.

6.7.7 Axis Definitions

X-axis — Feature occurrence number j : The j th time feature f is active within an episode. This is a within-episode counter that resets to 1 for each episode and is not a lesson-time index.

Y-axis — $\overline{\Delta P}_f^{(j)}$: Average change in productive-state probability from the episode opening, in probability units. A value of -0.18 means that at the j th occurrence of feature f , the IO-BKT posterior is on average 18 percentage points below where the episode started. Zero means the classroom is at its opening level. This axis is entirely in latent space; the raw annotation O_t does not appear.

6.7.8 Why $P(M_t=1)$ and Not O_t

The raw annotation $O_t \in \{0, 1\}$ is deliberately not used. First, O_t is a noisy emission: $O_t = 0$ can arise because $M_t = 0$ or because of annotation slip ($\hat{p}_S = 0.015$), and these cases are indistinguishable without the model. $P^{(e)}(M_t=1)$ integrates over the entire observation history and removes single-line noise. Second, O_t conflates regime with emission; $P(M_t=1)$ isolates the latent state. Third, $O_t \in \{0, 1\}$ is binary and would require additional smoothing to produce interpretable curves; the latent posterior is naturally continuous and probabilistically calibrated.

6.7.9 Comparison with Scarlatos et al.

Scarlatos et al. [1] plot:

$$\frac{1}{N} \sum_{i=1}^N \left(\hat{z}_{j,k}^{(i)} - \hat{z}_{1,k}^{(i)} \right) \quad (45)$$

where $\hat{z}_{j,k}^{(i)}$ is the predicted mastery of KC k at its j th occurrence in dialogue i , and the subtraction is from the *first occurrence of that KC* within that dialogue. This forces the curve to start at zero by construction. Our $\overline{\Delta P}_f^{(j)}$ subtracts the *episode opening* $P^{(e)}(M_1=1)$, not the first feature occurrence; consequently $\overline{\Delta P}_f^{(1)}$ is non-zero and reflects the productive-state level at the first occurrence of feature f relative to where the lesson started. For Behaviourism, $\overline{\Delta P}_{\text{Beh}}^{(1)} \approx -0.18$ means the classroom is already 18 percentage points below its opening baseline when the first Behaviourist line appears.

6.7.10 Interpretation of Fig. 3

Four patterns are present. **Stable suppression** (Behaviourism, $\overline{\Delta P} \approx -0.18$): the classroom sits persistently below its opening level at every Behaviourist occurrence; the flat trajectory confirms a consistent rather than transient association. **Stable near-zero** (Cognitivism, $\overline{\Delta P} \approx 0$): Cognitivist lines occur at roughly the episode’s own opening level. **Stable elevation** (Analysing, $\overline{\Delta P} \approx +0.20$; Constructivism, $\overline{\Delta P} \approx +0.15$): above-baseline association at every occurrence, consistent with positive estimates for these features. **Cumulative degradation** (Remembering, falling from $\overline{\Delta P}^{(1)} \approx -0.25$ to $\overline{\Delta P}^{(13)} \approx -0.40$): the only non-flat trajectory. Each additional Remembering occurrence




Figure 3: Knowledge Change Curves

is associated with a progressively lower productive-state probability, extending IO-BKT’s per-line $\hat{\beta}_{\text{off, Rem}}$ estimate to a longitudinal cumulative picture unavailable from regression or single-coefficient analysis.

Applying ($\overline{\Delta P} \approx 0$, stable) appears near baseline despite the largest ignition coefficient ($\hat{\beta}_{\text{on}} = +3.41$). Applying ignites the productive state when it is low and sustains it when it is already high; averaging across both contexts in Eq. (43) produces a near-zero mean, consistent with its symmetric presence in both and at nearly equal magnitudes (+3.41 and -3.40 respectively).

6.8 Teacher $\rightarrow$ Student Lag Analysis

6.8.1 Why “Lag” Analysis

In time-series analysis, a *lag- k* relationship measures the association between a variable at position t and another variable at position $t-k$. Here we identify for each student turn the *nearest preceding teacher turn* and record its EP code as the proximal instructional signal. Because teacher moves sometimes span multiple consecutive lines before a student responds, we search backwards through the episode until the nearest teacher line is found. This is consistent with the definition of a pedagogical move as a sustained teacher action rather than a single utterance.

6.8.2 Constructing Matched Pairs

The **nearest-teacher matched pair** at line t in episode e is:

$$(e, t) \in \mathcal{P} \iff \text{spk}_{e,t} = \text{Student} \wedge t' = \max\{s < t : \text{spk}_{e,s} = \text{Teacher}\} \quad (46)$$

That is, t' is the index of the most recent teacher line before student turn t , and the EP code of that line is the preceding move. The full matched-pair corpus and EP-stratified subsets are:

$$\mathcal{P}_c = \{(e, t) \in \mathcal{P} : \text{EP}_{e,t'} = c\}, \quad |\mathcal{P}_{\text{Beh}}|=1,582, |\mathcal{P}_{\text{Cog}}|=2,693, |\mathcal{P}_{\text{Con}}|=45 \quad (47)$$

giving $|\mathcal{P}|=4,320$ total matched pairs. Constructivist teacher moves are never the immediately adjacent line before a student response — they always appear with at least one further teacher line between the Constructivist utterance and the student turn. This is itself a structural finding: Constructivism in this corpus functions as sustained scaffolding embedded within teacher talk, not as a single-turn prompt-response exchange.

6.8.3 EP-Conditional Student 21CS Rate

The probability that a student demonstrates 21CS conditioned on the nearest preceding teacher EP is:

$$P(\text{student 21CS} \mid \text{EP}_{\text{preceding}} = c) = \frac{1}{|\mathcal{P}_c|} \sum_{(e,t) \in \mathcal{P}_c} O_{e,t} \quad (48)$$

where $O_{e,t} \in \{0,1\}$ is the raw 21CS annotation at the student turn, used directly without IO-BKT to provide annotation-level ground truth independent of any modelling assumption.

6.8.4 EP-Conditional Student Cognitive Process Level

The average Bloom's CP level of the student turn is:

$$\overline{\text{CP}}_c = \frac{1}{|\mathcal{P}_c|} \sum_{(e,t) \in \mathcal{P}_c} \text{CP}_{e,t} \quad (49)$$

where $\text{CP}_{e,t} \in \{1, \dots, 6\}$ is the Bloom's level of the student utterance (1 = Remembering, ..., 6 = Creating).

6.8.5 Lesson-Phase Analysis via Global Encoding

Rather than partitioning lessons into arbitrary equal-width thirds, we use the corpus's **Global Encoding** (GE) variable, which assigns each line to one of three pedagogically grounded phases: $g \in \{\text{Preassessment}, \text{Teaching}, \text{Assessment}\}$. The GE column uses sparse encoding — only the first row of each phase is marked; remaining rows are propagated via within-episode forward-fill. For each phase g and EP category c :

Table 2: Nearest-teacher lag analysis: teacher EP $\rightarrow$ student 21CS.

Preceding EP	$P(\text{student 21CS})$	Mean CP	n
Behaviourism	0.487	2.17	1,582
Cognitivism	0.536	2.12	2,693
Constructivism	0.978	2.12	45
<i>By Global Encoding phase (Eq. 50):</i>			
Preassessment $\times$ Behaviourism	0.100	–	194
Preassessment $\times$ Cognitivism	0.220	–	290
Teaching $\times$ Behaviourism	0.390	–	848
Teaching $\times$ Cognitivism	0.490	–	1,935
Teaching $\times$ Constructivism	1.000	–	1
Assessment $\times$ Behaviourism	0.779	–	540
Assessment $\times$ Constructivism	0.955	–	44

Notes: CP level: 1 = Remembering, ..., 6 = Creating. Overall student 21CS prevalence across all matched pairs = 0.536. Constructivism never appears as the strictly adjacent teacher turn (lag=1); all 45 pairs are found at lag>1, reflecting its role as sustained scaffolding rather than a single-turn prompt. Assessment $\times$ Cognitivism omitted ($n<5$ in this corpus version).

$$P_g(\text{student 21CS} \mid c) = \frac{\sum_{\substack{(e,t) \in \mathcal{P}_c \\ \text{GE}_{e,t'}=g}} O_{e,t}}{|\{(e,t) \in \mathcal{P}_c : \text{GE}_{e,t'} = g\}|} \quad (50)$$

where GE is read from the teacher line t' , not the student line. For cross-corpus comparability (GE may not be available in other datasets), we also report the lesson-thirds partition:

$$P_k(\text{student 21CS} \mid c) = \frac{\sum_{\substack{(e,t) \in \mathcal{P}_c \\ \tau_{e,t} \in \mathcal{I}_k}} O_{e,t}}{|\{(e,t) \in \mathcal{P}_c : \tau_{e,t} \in \mathcal{I}_k\}|}, \quad \mathcal{I}_k = \begin{cases} [0, \frac{1}{3}) & k=1 \\ [\frac{1}{3}, \frac{2}{3}) & k=2 \\ [\frac{2}{3}, 1] & k=3 \end{cases} \quad (51)$$

6.8.6 Results

Table 2 and Fig. 4 report the values of Eqs. (48)–(50).

Overall EP effect (Panel a). A Constructivist teacher move is followed by student 21CS in 97.8% of cases (Eq. 48, $c=\text{Con}$, $n=45$), nearly double the Behaviourist rate (48.7%) and substantially above the Cognitivist rate (53.6%). The overall student 21CS prevalence across all matched pairs is 0.536.

Student cognitive depth (Panel b). Average student CP level is near-identical across all three EP conditions ($\overline{\text{CP}} \approx 2.12\text{--}2.17$, corresponding to Understanding-level responses, Eq. 49). The markedly higher 21CS rate following Constructivist moves is therefore not explained by deeper cognitive responses but by more frequent activation of the productive-discourse regime itself.




Figure 4: Teacher → Student Lag Analysis (Eqs. 48–50, $|\mathcal{P}|=4,320$ matched pairs across 48 episodes). Panel (a): $P(\text{student 21CS} \mid \text{EP}_{\text{preceding}})$ from Eq. (48); Constructivist moves ($n=45$) are followed by student 21CS in 97.8% of cases. Panel (b): average student Bloom’s CP level (Eq. 49); CP is near-uniform across EP conditions (≈ 2.12 , Understanding level), ruling out cognitive depth as an explanatory variable. Panel (c): EP effect across lesson thirds (Eq. 51); Behaviourism shows a monotonic increase across lesson position ($0.407 \rightarrow 0.469 \rightarrow 0.591$). Panel (d): EP effect by Global Encoding phase (Eq. 50); all EPs produce near-zero 21CS rates in Preassessment and high rates in Assessment, with Behaviourism rising from $P=0.100$ to $P=0.779$, demonstrating lesson phase as a strong moderator of the EP-21CS association.

Lesson phase via Global Encoding (Panel d). The GE analysis reveals the strongest moderating pattern in the corpus. The same Behaviourist teacher move produces $P=0.100$ in Preassessment but $P=0.779$ in Assessment, a nearly eightfold increase. Cognitivism rises from 0.220 to a similarly high Assessment-phase rate. This reflects the pedagogical function of each phase: Preassessment elicits prior knowledge before instruction (most students cannot yet demonstrate 21CS), while Assessment probes consolidated understanding (students who have engaged with the lesson can demonstrate it regardless of the specific EP used to elicit the response). Constructivism concentrates in the Assessment phase ($n=44$, $P=0.955$), with only a single occurrence in Teaching ($P=1.000$, $n=1$), indicating that this corpus’s rare Constructivist moves are deployed primarily as assessment-type scaffolding rather than as within-instruction discovery tasks.

These GE findings establish that the aggregate EP effect in Panel (a) is substantially moderated by lesson phase, and that Behaviourism is not uniformly suppressive: in the Assessment phase, Behaviourist moves elicit high 21CS rates because the productive regime has already been established by prior instruction.

Lesson thirds (Panel c). The lesson-thirds partition (Eq. 51) shows Constructivist moves achieving $P \approx 1.0$ across all three thirds, consistent with the Assessment-phase GE finding. Behaviourism shows a modest monotonic increase (Opening $0.407 \rightarrow$ Middle $0.469 \rightarrow$ Closing 0.591), while Cognitivism remains near 0.52 throughout, confirming that its association with 21CS is stable across lesson position.

Link to IO-BKT. The quantities in Eq. (48) are computed entirely from raw annotations without model assumptions and therefore constitute observational ground truth for IO-BKT’s transition coefficients. The GE finding adds a layer the coefficient table cannot show: $\hat{\beta}_{\text{on,Con}} = +1.25$ captures the average ignition effect of Constructivist moves; the GE analysis reveals that this effect is concentrated in Assessment-phase interactions ($n=44$, $P=0.955$), where Constructivist scaffolding meets students who have had the full lesson

to develop their understanding. The per-student tracking of what follows, whether the individual student has mastered the underlying knowledge component, is the task that student-level KT [1] would address downstream.

References

- [1] A. Scarlatos, R. S. Baker, and A. Lan, “Exploring knowledge tracing in tutor-student dialogues using LLMs,” in *Proc. LAK*, 2025, pp. 249–259.